

Discriminative Classifiers

Goals

- Introduce 2 new classifiers
- Understand how to fit them
- Think about their drawbacks
- Contrast with logistic regression
- Just an intro, not an exhaustive list(!), notable omissions:
 - SVMs
 - Generative classifiers


What are Discriminative Classifiers?

- Learn a function that directly maps from an input to a class


What are Discriminative Classifiers?

- Learn a function that directly maps from an input to a class
- This is equivalent to learning a **boundary** separating classes


What are Discriminative Classifiers?




K-Nearest-Neighbours




K-Nearest-Neighbours



K-Nearest-Neighbours



K-Nearest-Neighbours



K-Nearest-Neighbours: Comments (1)

- Alternatively, neighbourhoods may be defined by a radius

K-Nearest-Neighbours: Comments (1)

- Alternatively, neighbourhoods may be defined by a radius
- In the binary classification case, keep K odd to avoid a tied vote

K-Nearest-Neighbours: Comments (1)

- Alternatively, neighbourhoods may be defined by a radius
- In the binary classification case, keep K odd to avoid a tied vote
- KNN is an example of **instance-based learning**

K-Nearest-Neighbours: Comments (2)

- The computational cost resides in the assignment of a class to a new point (req. computing the distance to all training points):

$$O(TN(D + K))$$

- **T**: number of test points to classify
- **N**: number of training points
- **D**: complexity associated with computing the distance between two points
- **K**: parameter of the KNN

Visualising KNN Decision Boundaries


- Rule of thumb:
 - **K↓** boundaries rougher → more local (risk overfitting)
 - **K↑** boundaries smoother → more global (risk underfitting)

Visualising KNN Decision Boundaries


- Rule of thumb:
 - **K↓** boundaries rougher → more local (risk overfitting)
 - **K↑** boundaries smoother → more global (risk underfitting)

So... how can we pick a good K?


Overfitting & Underfitting




Overfitting & Underfitting




Overfitting & Underfitting




Bias & Variance



Bias & Variance



Bias & Variance

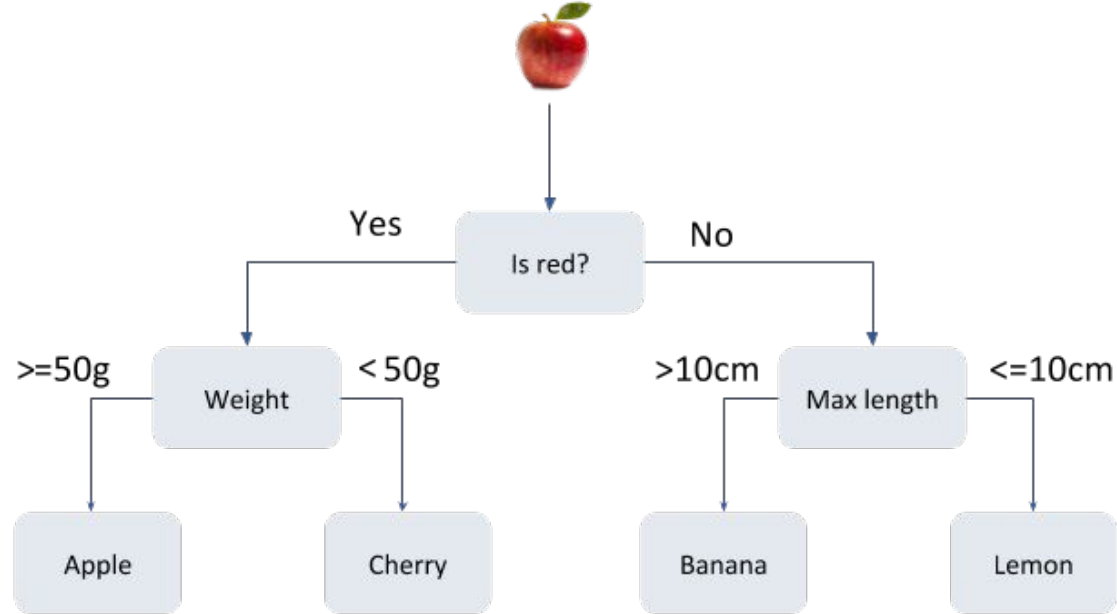




Hands-on session

knn-dtc.ipynb
(first half - 30 mins)

Decision Tree Classifiers (DTC)



Decision Tree Classifiers (DTC)

- A DTC is a model with a **hierarchical** structure. Each node corresponds to a feature and a split. To fit a DTC:

Decision Tree Classifiers (DTC)

- A DTC is a model with a **hierarchical** structure. Each node corresponds to a feature and a split. To fit a DTC:
 - At each level, consider all features and for each, the splitting point that best separates the classes

Decision Tree Classifiers (DTC)

- A DTC is a model with a **hierarchical** structure. Each node corresponds to a feature and a split. To fit a DTC:
 - At each level, consider all features and for each, the splitting point that best separates the classes
 - Take the feature that offers the best separation and continue

Decision Tree Classifiers (DTC)

- But what is **separation**?

Decision Tree Classifiers (DTC)

- But what is **separation**?
 - Different metrics can be used to measure **separation**

Decision Tree Classifiers (DTC)

- But what is **separation**?
 - Different metrics can be used to measure **separation**
 - Common ones are **GINI impurity** and **information gain**

GINI Impurity

- Assume there are k classes: $1, 2, \dots, k$
- At a given node, P , there is a proportion p_1 of elements in class 1, p_1 of elements in class 2, etc

GINI Impurity

- Assume there are k classes: $1, 2, \dots, k$
- At a given node, P , there is a proportion p_1 of elements in class 1, p_2 of elements in class 2, etc
- GINI Impurity at that node:

$$G(P) = 1 - \sum_{i=1}^k p_i^2$$

Fitting a tree with GINI Impurity

- At each node, the parent node P must be split into children nodes c_1, \dots, c_J
- Each child node has n_j instances

Fitting a tree with GINI Impurity

- At each node, the parent node P must be split into children nodes c_1, \dots, c_J
- Each child node has n_j instances
- So, the **Gain** for the split:

$$Gain = G(P) - \frac{\sum_j n_j G(c_j)}{\sum_j n_j} \sum_{i=1}^k p_i^2$$


When to stop splitting?

- Pick a **maximum depth**

When to stop splitting?

- Pick a **maximum depth**
- Otherwise, we get one sample per “leaf node” → overfitting

Decision Tree Classifiers (DTC)





Hands-on session

knn-dtc.ipynb
(second half - 30 mins)